


Efficient-Parallel-C++

General Weighted Matching – Alan Mazankiewicz

Institute for Theoretical Informatics – Prof. Sanders



Overview



Matching Problem



Algorithms


- Local Dominant
- Suitor
- Heavy Matching



Evaluation


Definition - Matching

A **matching** in an undirected Graph $G = (V, E)$ is a subset of edges M such that no two elements of M have common endpoints




Definition - Matching

A **matching** in an undirected Graph $G = (V, E)$ is a subset of edges M that such no two elements of M have common endpoints




Definition – Maximum Weighted Matching

A **maximum weighted matching** of an undirected graph $G = (V, E, W)$ is a matching M with the largest possible sum* of weights



Definition – Maximum Weighted Matching

A **maximum weighted matching** of an undirected graph $G = (V, E, W)$ is a matching M with the largest possible sum* of weights



$$O(nm + n^2 \log n) \text{ [Gabow 90]}$$



Half-Approximation Algorithms

Sequential

- Greedy – Avis
- PGA' - Path Growing Algorithm - Drake and Hougardy
- GPA - Global Paths Algorithm - Maue and Sanders
- HEM – Heuristics of Heavy Matching – Birn et al.

Sequential and Parallel

- Locally Dominant Edges – Halappanavar et al.
- Suitor – Manne and Halappanavar



Half-Approximation Algorithms

Sequential

- Greedy – Avis
- PGA' - Path Growing Algorithm - Drake and Hougardy
- GPA - Global Paths Algorithm - Maue and Sanders
- HEM – Heuristics of Heavy Matching – Birn et al.


Sequential and Parallel

- **Locally Dominant Edges** – Halappanavar et al.
- **Suitor** – Manne and Halappanavar




Greedy


- Sort edges in descending order




Greedy



Greedy



Greedy




Greedy



$$O(m + m \log n)$$




Local Dominant - Sequential




Local Dominant - Sequential

- Phase I
- Process each Vertex




Local Dominant - Sequential

- Phase I
- Process each Vertex




Local Dominant - Sequential

- Phase I
- Process each Vertex




Local Dominant - Sequential

- Phase I
- Add Endpoints to Queue Q




Local Dominant - Sequential

- Phase II - *while* $Q \neq \{\}$




Local Dominant - Sequential

- Phase II - *while* $Q \neq \{\}$
- Process each vertex that endpoint is candidate of




Local Dominant - Sequential

- Phase II - *while* $Q \neq \{\}$
- Process each vertex that endpoint is candidate of




Local Dominant - Sequential

- Phase II - *while* $Q \neq \{\}$
- Process each vertex that endpoint is candidate of



Local Dominant - Sequential

- Same matching as Greedy



$$O(n + m * \Delta)$$



Local Dominant - Parallel

```

1: procedure PARALLEL-QUEUE( $G(V, E)$ ,
   mate)
2:    $Q_C \leftarrow \emptyset$ 
3:    $Q_N \leftarrow \emptyset$ 
4:   — — Phase-1 — —
5:   for each  $u \in V$  in parallel do

6:     mate[ $u$ ]  $\leftarrow \emptyset$ 
7:     candidate[ $u$ ]  $\leftarrow \emptyset$ 
8:     max_wt  $\leftarrow -\infty$ 
9:     max_wt_id  $\leftarrow \emptyset$ 
10:    for each  $v \in \text{adj}(u)$  do
11:      if (mate[ $v$ ] =  $\emptyset$ ) AND (max_wt
        <  $w(e_{u,v})$ ) then
12:        max_wt  $\leftarrow w(e_{u,v})$ 
13:        max_wt_id  $\leftarrow v$ 
14:      candidate[ $u$ ]  $\leftarrow$  max_wt_id
15:    for each  $u \in V$  in parallel do

16:      if candidate[candidate[ $u$ ]] =  $u$ 
        then
17:        mate[ $u$ ]  $\leftarrow$  candidate[ $u$ ]
18:         $Q_C \leftarrow Q_C \cup \{u\}$ 

19:
20:   — — Phase-2 — —
21:   while  $Q_C \neq \emptyset$  do
22:     for each  $u \in Q_C$  in parallel
       do
23:       for each  $v \in \text{adj}(u) \setminus$  do
24:         if candidate[ $v$ ] =  $u$  then

25:           PROCESSVERTEX( $v, Q_N$ )
26:            $Q_C \leftarrow Q_N$ 
27:            $Q_N \leftarrow \emptyset$ 

```



Source: Halappanavar et al.

Local Dominant - Parallel

```

void ParallelDominant::processVertex(NodeHandle u, moodycamel::ConcurrentQueue<NodeHandle> &Q, Mate &candidate)
{
    double max_wt = 0;
    NodeHandle max_wt_id = G.null_node();
    size_t u_id = G.nodeId(u);

    for (EdgeIterator t = G.beginEdges(u); t != G.endEdges(u); ++t)
    {
        if ((mate[*t] == G.null_node()) && (max_wt < G.edgeWeight(t)))
        {
            max_wt = G.edgeWeight(t);
            max_wt_id = G.node(*t);
        }
    }
    candidate[u_id] = max_wt_id;

    if (candidate[u_id] == G.null_node())
        return;

    if (G.nodeId(candidate[G.nodeId(candidate[u_id])]) == u_id)
    {
        mate[u_id] = candidate[u_id];
        mate[G.nodeId(candidate[u_id])] = u;
        Q.enqueue(u);
        Q.enqueue(candidate[u_id]);
    }
}

class ParallelDominant
{
public:
    using EdgeIterator = AdjacencyArray::EdgeIterator;
    using NodeHandle = AdjacencyArray::NodeHandle;
    using Mate = std::vector<NodeHandle>;

```



Suitor – Sequential

■ *suitor*(*u*) - Who points to *u*

```

1: for each  $u \in V$  do
2:   suitor(u) = NULL
3:   ws(u) = 0
4: for each  $u \in V$  do
5:   current = u
6:   done = False
7:   while not done do
8:     partner = suitor(current)
9:     heaviest = ws(current)
10:    for each  $v \in N(\textit{current})$  do
11:      if  $w(\textit{current}, v) >$  heaviest and
         $w(\textit{current}, v) > ws(v)$  then
12:        partner = v
13:        heaviest =  $w(\textit{current}, v)$ 
14:    done = True
15:    if heaviest > 0 then
16:      y = suitor(partner)
17:      suitor(partner) = current
18:      ws(partner) = heaviest
19:      if y ≠ NULL then
20:        current = y
21:        done = False

```



Source: Manne and Halappanavar

Suitor – Sequential

■ $O(m * \Delta)$

```

1: for each  $u \in V$  do
2:    $suitor(u) = NULL$ 
3:    $ws(u) = 0$ 
4: for each  $u \in V$  do
5:    $current = u$ 
6:    $done = False$ 
7:   while not  $done$  do
8:      $partner = suitor(current)$ 
9:      $heaviest = ws(current)$ 
10:    for each  $v \in N(current)$  do
11:      if  $w(current, v) > heaviest$  and
         $w(current, v) > ws(v)$  then
12:         $partner = v$ 
13:         $heaviest = w(current, v)$ 
14:     $done = True$ 
15:    if  $heaviest > 0$  then
16:       $y = suitor(partner)$ 
17:       $suitor(partner) = current$ 
18:       $ws(partner) = heaviest$ 
19:      if  $y \neq NULL$  then
20:         $current = y$ 
21:         $done = False$ 

```



Source: Manne and Halappanavar

Suitor – Parallel with Lock

```

void ParallelSuitor::process(NodeHandle start, NodeHandle finish)
{
    for (NodeHandle u = start; u != finish; ++u)
    {
        NodeHandle current = u;
        bool done = false;

        while (!done)
        {
            NodeHandle partner = G.null_node();
            double heaviest = 0;

            for (EdgeIterator v = G.beginEdges(current); v != G.endEdges(current); ++v)
            {
                double weight_v = G.edgeWeight(v);
                if (weight_v > heaviest && weight_v > ws[*v])
                {
                    partner = G.node(*v);
                    heaviest = weight_v;
                }
            }
            done = true;
            if (heaviest > 0)
            {
                std::unique_lock<std::mutex> lock(mtx[G.nodeId(partner)]);

                if (ws[G.nodeId(partner)] < heaviest)
                {
                    NodeHandle y = suitor[G.nodeId(partner)];
                    suitor[G.nodeId(partner)] = current;
                    ws[G.nodeId(partner)] = heaviest;
                    lock.unlock();

                    if (y != G.null_node())
                    {
                        current = y;
                        done = false;
                    }
                } else
                {
                    done = false;
                }
            }
        }
    }
}

```



Graph Data Structure

- Nodes: [0, 3, ...] → NodeHandle
- Edges: [1,2,3, 4,5,6 ...] → Edgelterator
- Weights: [11,12,13, 14, 15, 16, ...] → Index of Edgelterator

Suitor – Parallel Lockfree

```

void ParallelLocklessSuitor::process(NodeHandle start, NodeHandle finish)
{
    for (NodeHandle u = start; u != finish; ++u)
    {
        NodeHandle current = u;
        bool done = false;

        while (!done)
        {
            EdgeIterator partner = G.null_edge();
            double heaviest = 0;

            for (EdgeIterator v = G.beginEdges(current); v != G.endEdges(current); ++v)
            {
                double weight_v = v->second;
                if (weight_v > heaviest && weight_v > suitor[v->first].load()->second)
                {
                    partner = v;
                    heaviest = weight_v;
                }
            }
            done = true;
            if (heaviest > 0)
            {
                size_t partnerId = partner->first;
                EdgeIterator current_suitor = suitor[partnerId];
                if (current_suitor->second < heaviest)
                {
                    if (suitor[partnerId].compare_exchange_strong(current_suitor, partner))
                    {
                        NodeHandle y = G.edgeHead(current_suitor);
                        if (y != G.null_node())
                        {
                            current = y;
                            done = false;
                        }
                    } else
                    {
                        done = false;
                    }
                } else
                {
                    done = false;
                }
            }
        }
    }
}

```



Graph Data Structure

- Nodes: [0, 3, ...] → NodeHandle
- Edges: [1,2,3, 4,5,6 ...] → Edgelterator
- Weights: [11,12,13, 14, 15, 16, ...] → Index of Edgelterator

- Nodes: [0, 3, ...] → NodeHandle
- Edges: [(1,11),(2, 12),(3,13), (4, 14) ...] → Edgelterator

Suitor – Parallel Lockfree

```

void ParallelLocklessSuitor::process(NodeHandle start, NodeHandle finish)
{
    for (NodeHandle u = start; u != finish; ++u)
    {
        NodeHandle current = u;
        bool done = false;


        while (!done)
        {
            EdgeIterator partner = G.null_edge();
            double heaviest = 0;

            for (EdgeIterator v = G.beginEdges(current); v != G.endEdges(current); ++v)
            {
                double weight_v = v->second;
                if (weight_v > heaviest && weight_v > suitor[v->first].load()->second)
                {
                    partner = v;
                    heaviest = weight_v;
                }
            }
            done = true;
            if (heaviest > 0)
            {
                size_t partnerId = partner->first;
                EdgeIterator current_suitor = suitor[partnerId];
                if (current_suitor->second < heaviest)
                {
                    if (suitor[partnerId].compare_exchange_strong(current_suitor, partner))
                    {
                        NodeHandle y = G.edgeHead(current_suitor);
                        if (y != G.null_node())
                        {
                            current = y;
                            done = false;
                        }
                    } else
                    {
                        done = false;
                    }
                } else
                {
                    done = false;
                }
            }
        }
    }
}

```




Runtime optimized vs unoptimized




Quality Improvement – Heavy Matching

■ 1st Matching




Quality Improvement – Heavy Matching

■ 2nd Matching




Quality Improvement – Heavy Matching


- Path or Cycle
- Can be parallelized




Runtime Sequential Implementation



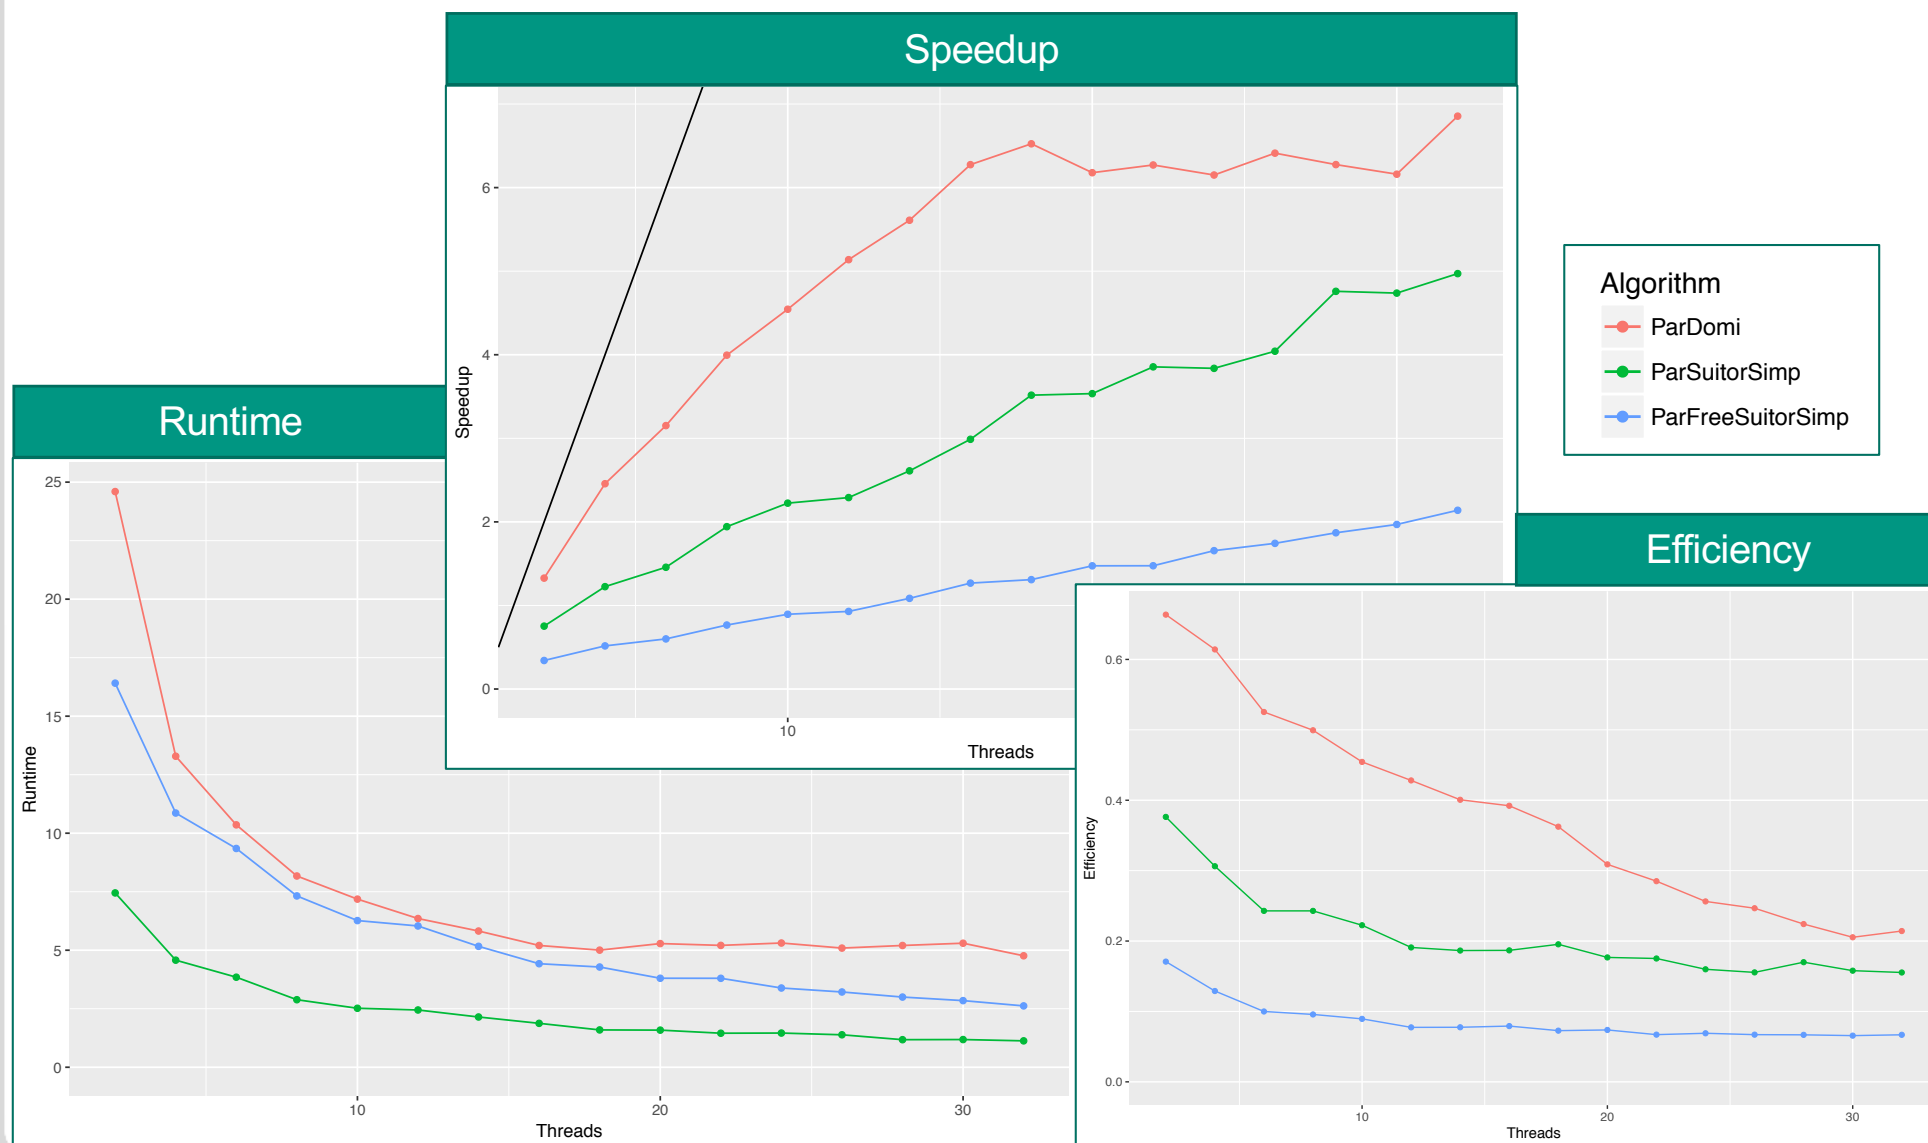
Scalability Sequential Implementation



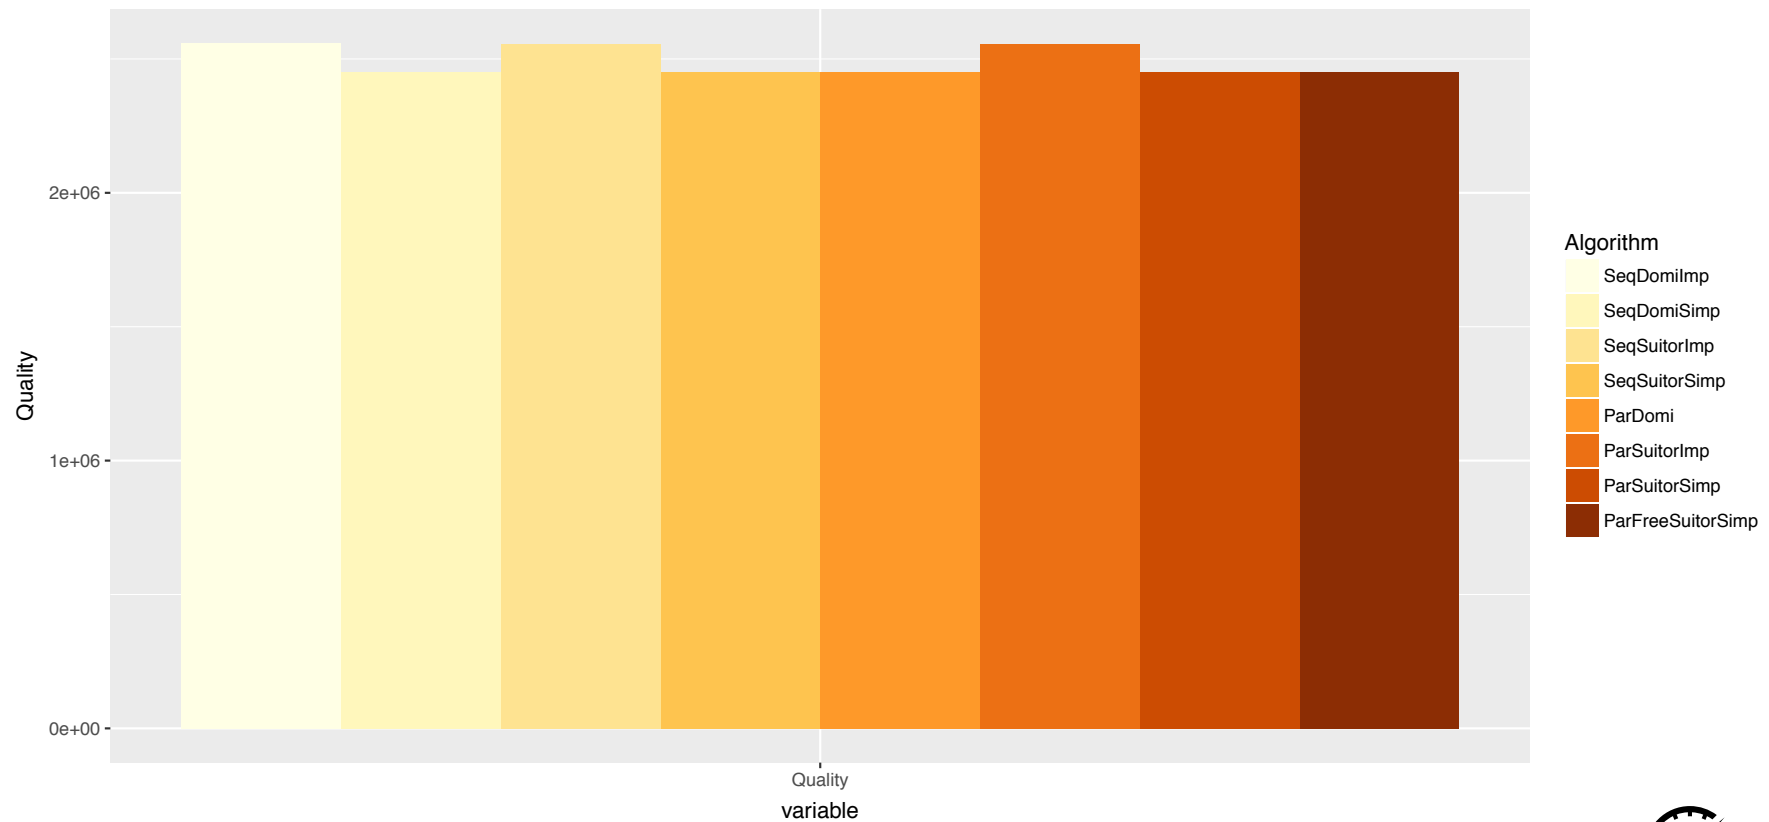
Runtime Heavy vs Simple Matching



Evaluation



Quality Comparison



Summary



Matching Problem



Algorithms

- Local Dominant
- Suitor
- Heavy Matching



Evaluation


References

- J. Maue and P. Sanders, “Engineering algorithms for approximate weighted matching,” in *WEA’07*. LNCS, Springer, 2007, pp. 242–255.
- M. Halappanavar, J. Feo, O. Villa, A. Tumeo, and A. Pothen, “Approximate weighted matching on emerging manycore and multithreaded architectures,” *Int. J. High Perf. Comput. App.*, vol. 26, no. 4, pp. 413– 430, 2012.
- Manne, F., & Halappanavar, M. (2014, May). New effective multithreaded matching algorithms. In *2014 IEEE 28th International Parallel and Distributed Processing Symposium*(pp. 519-528). IEEE.

Backup


Definition – Maximal Matching

A **maximal matching** is a matching M of an undirected graph $G = (V, E)$ with the property that if any edge not in M is added to M , it is no longer a matching




Definition – Maximal Matching

A **maximal matching** is a matching M of an undirected graph $G = (V, E)$ with the property that if any edge not in M is added to M , it is no longer a matching




Definition – Maximal Matching

A **maximal matching** is a matching M of an undirected graph $G = (V, E)$ with the property that if any edge not in M is added to M , it is no longer a matching




Definition – Maximum Cardinality Matching

A **maximum cardinality matching** of an undirected graph $G = (V, E)$ is a matching M that contains the largest possible number of elements



Backup - Evaluation



*Parallel with 32 threads

Backup – Complexities

- Greedy: $O(m + m \log n)$
- PGA': $O(m)$
- Local Dominant: $O(n + m * \Delta)$
- Suitor: $O(m * \Delta)$
- Gabow: $O(nm + n^2 \log n)$


Backup – Graphs

- Runtime Sequential Implementation
 - nVertices 1000000
 - nEdges 100000000
- Runtime Heavy vs Simple Matching & Quality Comparison
 - nVertices 100000
 - nEdges 1000000
- Parallel Running Time
 - nVertices 6000000
 - nEdges 400000000

Backup - Graphs

■ Scalability Sequential Implementation

- Sec. 1: nVertices 100.000 nEdges 1.000.000
- Sec. 2: nVertices 200000 nEdges 2.000.000
- Sec. 3: nVertices 600.000 nEdges 6.000.000
- Sec. 4: nVertices 1.000.000 nEdges 1.000.000.000
- Sec. 5: nVertices 2.000.000 nEdges 2.000.000.000
- Sec. 6: nVertices 4.000.000 nEdges 4.000.000.000



Backup – Sequential Local Dominant

```

1: procedure SERIAL-QUEUE
   ( $G(V, E)$ ,  $\text{mate}$ )
2: for each  $u \in V$  do
3:    $\text{mate}[u] \leftarrow \emptyset$ 
4:    $\text{candidate}[u] \leftarrow \emptyset$ 
5:    $Q \leftarrow \emptyset$ 
6:   for each  $u \in V$  do
7:      $\text{PROCESS\_VERTEX}(u, Q)$ 
8:   while  $Q \neq \emptyset$  do
9:      $u \leftarrow \text{FRONT}(Q)$ 
10:     $Q \leftarrow Q \setminus \{u\}$ 
11:    for each  $v \in \text{adj}(u) \setminus \text{mate}(u)$  do
12:      if  $\text{candidate}[v] = u$  then
13:         $\text{PROCESS\_VERTEX}(v, Q)$ 

```

```

1: procedure PROCESSVERTEX ( $s, Q$ )
2:    $\text{max\_wt} \leftarrow -\infty$ 
3:    $\text{max\_wt\_id} \leftarrow \emptyset$ 
4:   for each  $t \in \text{adj}(s)$  do
5:     if ( $\text{mate}[t] = \emptyset$ ) AND
       ( $\text{max\_wt} < w(e_{s,t})$ ) then
6:        $\text{max\_wt} \leftarrow w(e_{s,t})$ 
7:        $\text{max\_wt\_id} \leftarrow t$ 
8:    $\text{candidate}[s] \leftarrow \text{max\_wt\_id}$ 
9:   if  $\text{candidate}[\text{candidate}[s]] = s$ 
     then
10:     $\text{mate}[s] \leftarrow \text{candidate}[s]$ 
11:     $\text{mate}[\text{candidate}[s]] \leftarrow s$ 
12:     $Q \leftarrow Q \cup \{s, \text{candidate}[s]\}$ 

```

Source: Halappanavar et al.

Backup – Evaluation Suitor

	<i>Fl. Small</i>	<i>Fl. Large</i>	<i>Rand</i>	<i>RMAT</i>	<i>Geo.</i>	<i>Comp. Geo</i>	<i>Comp. Rand</i>
GREEDY	56.1	56.7	31.4	39.2	43.8	108.7	105.8
GREEDY-	1.0	1.5	0.8	1.3	1.8	1.7	1.3
PGA'	3.2	4.5	4.2	5.5	4.9	2.7	2.7
LOCALMAX	8.2	5.5	5.6	9.9	4.0	1166.0	5.8
SUITOR	2.5	2.8	3.1	4.1	3.1	195.6	2.4
SORTSUITOR	42.7	31.7	15.9	16.8	16.9	118.1	116.1
SORTSUITOR-	1.3	1.5	1.7	2.4	1.9	0.9	0.01

TABLE I
TIME TAKEN RELATIVE TO THE HEM ALGORITHM.

	<i>Fl. Small</i>	<i>Fl. Large</i>	<i>Rand</i>	<i>RMAT</i>	<i>Geo</i>	<i>Comp. Geo</i>	<i>Comp. Rand</i>
GPA	22.8	21.9	12.3	11.3	15.7	55.7	54.2
GPA-	0.9	2.3	1.4	1.8	3.5	1.1	0.8
2RPGA'	3.2	3.5	3.1	3.2	3.5	3.1	3.1
2RSUITOR	2.1	2.0	2.6	2.6	2.2	199.2	2.5
2RSORTSUITOR	18.0	13.0	6.6	5.4	5.9	60.7	59.3
2RSORTSUITOR-	1.2	1.6	1.4	1.8	1.5	0.9	0.01


TABLE II
TIME TAKEN RELATIVE TO THE 2RHEM ALGORITHM.

	<i>Fl. Small</i>	<i>Fl. Large</i>	<i>Rand</i>	<i>RMAT</i>	<i>Geo</i>	<i>Comp. Geo</i>	<i>Comp. Rand</i>
2RSUITOR	0.9986	0.9977	0.9983	1.0027	0.9989	1.0000	0.9995
2RPGA'	0.9972	0.9966	0.9960	0.9983	0.9969	0.9994	0.9997
2RHEM	0.8172	0.7795	0.9762	0.9675	1.0007	0.9959	0.9995
GREEDY	0.9921	0.9790	0.9761	0.9626	0.9749	0.9999	0.9992


TABLE III
RELATIVE WEIGHT OF TWO-ROUND ALGORITHMS COMPARED TO GPA.

Source: Manne and Halappanavar

Backup – Evaluation Suitor



(c) SUITOR on *RMAT* and *Geometric* graphs



(d) 2RSUITOR on *RMAT* and *Geometric* graphs

Source: Manne and Halappanavar

Suitor – Parallel Lockfree

Lock

```
double weight_v = G.edgeWeight(v);  
if (weight_v > heaviest && weight_v > ws[*v])
```

Lockfree - unoptimized


```
double weight_v = G.edgeWeight(v);  
if (weight_v > heaviest && weight_v > G.edgeWeight(suitor[*v]))
```

Lockfree - optimized


```
double weight_v = v->second; |  
if (weight_v > heaviest && weight_v > suitor[v->first].load()->second)
```



Quality Improvement – Heavy Matching




Quality Improvement – Heavy Matching





Quality Improvement – Heavy Matching


- Extend to maximal matching



Backup – Quality Scaled




Quality Improvement – Heavy Matching



Quality Improvement – Heavy Matching

■ Dynamic Programming



Backup – Greedy Algorithm

```
GRDY( $G = (V, E)$ ,  $w : E \rightarrow \mathbb{R}^{\geq 0}$ )  
1  $M := \emptyset$   
2 while  $E \neq \emptyset$  do  
3   let  $e$  be the edge with biggest weight in  $E$   
4   add  $e$  to  $M$   
5   remove  $e$  and all edges adjacent to its endpoints from  $E$   
6 return  $M$ 
```

Fig. 1. The greedy algorithm for approximate weighted matchings.

Backup – PGA' Algorithm

```

PGA'( $G = (V, E)$ ,  $w: E \rightarrow \mathbb{R}^{\geq 0}$ )
1   $M := \emptyset$ 
2  while  $E \neq \emptyset$  do
3     $P := \langle \rangle$ 
4    arbitrarily choose  $v \in V$  with  $\deg(v) > 0$ 
5    while  $\deg(v) > 0$  do
6      let  $e = (v, u)$  be the heaviest edge adjacent to  $v$ 
7      append  $e$  to  $P$ 
8      remove  $v$  and its adjacent edges from  $G$ 
9       $v := u$ 
10    $M := M \cup \text{MaxWeightMatching}(P)$ 
11   extend  $M$  to a maximal matching
12  return  $M$ 

```

Fig. 2. The improved Path Growing Algorithm PGA'.


Backup – Dynamic Programming


```


MaxWeightMatching( $P = \langle e_1, \dots, e_k \rangle$ )
1   $W[0] := 0; \quad W[1] := w(e_1)$ 
2   $M[0] := \emptyset; \quad M[1] := \{e_1\}$ 
3  for  $i := 2$  to  $k$  do
4    if  $w(e_i) + W[i - 2] > W[i - 1]$  then
5       $W[i] := w(e_i) + W[i - 2]$ 
6       $M[i] := M[i - 2] \cup \{e_i\}$ 
7    else
8       $W[i] := W[i - 1]$ 
9       $M[i] := M[i - 1]$ 
10 return  $M[k]$ 


```


Fig. 7. Obtaining a maximum weight matching for a path by dynamic programming.


























$O(m)$








- Extend to maximal matching





■ Dynamic Programming

